

Qualified Tender Change Intelligence

A procurement-data pipeline that separates republication noise from material changes, then routes each result through explainable qualification rules.

The operational problem

An update label is not proof that a bid-relevant field changed. A small pursuit team needs fewer false alarms, clear rejection reasons, and a safe review path when the source lacks evidence.

6

update-tagged
comparisons

0

material changes
found

1

real cancellation
surfaced

7/7

automated tests
passing

LIVE CASE A | NOISE SUPPRESSION

One procurement process contained 7 releases. Six later releases carried **tenderUpdate**, yet none changed status, deadline, value, currency, classification, or lot state. All six were rejected as **NO_MATERIAL_CHANGE**.

LIVE CASE B | SAFE ESCALATION

A second process changed from active to complete, removed its deadline and value, and closed its lot. The pipeline surfaced 5 field changes, then routed the record to **review-needed** because qualification evidence was absent.

What this demonstrates

- Direct field comparison instead of trusting source labels
- Reason codes and provenance on every retained decision
- Deterministic accepted, rejected, and review-needed queues
- No persisted contact details, descriptions, or unstructured prose

From public record to decision queue

A small, auditable pipeline: one sanctioned API path, a strict persistence boundary, deterministic transformations, and human review where evidence is incomplete.



Declared material field set

`tender.status`

`tender.value.amount`

`tender.classification.id`

`tender.tenderPeriod.endDate`

`tender.value.currency`

`tender.lots`

Qualification is a separate decision

A material change is not automatically a good lead. The profile independently checks classification prefix, value range, and deadline threshold. Failed rules are rejected; missing inputs go to review.

This separation prevents a change detector from silently becoming a lead-scoring black box.

Evidence retained with every result

- source URL + retrieval timestamp
- SHA-256 source fingerprint
- release ID + procurement process ID
- changed fields + old/new values
- decision + reason codes

What was verified - and what was not claimed

The portfolio value is operational honesty: reproducible evidence, visible failure paths, and strict limits around source access and personal data.

Verification matrix

CONTROL	EVIDENCE	OUTCOME
Update tag without declared field change	6 live comparisons	Rejected as noise
Material cancellation/status transition	1 live comparison	Five changes surfaced
Missing qualification evidence	1 live comparison	Sent to review-needed
Sanitized capture replay	Automated test	Same result after reload
Contact and prose exclusion	Automated test	Absent from saved records
Fixture behavior suite	7 automated tests	All passing

COMPLIANCE BY DESIGN

Two explicit record-package requests. No login, browser scraping, pagination, retry, proxying, rate-limit discovery, or adaptive polling. Raw responses remained in memory and were reduced to an allowlist before disk persistence.

LIMITATIONS STATED UP FRONT

This sample does not prove source-wide change frequency, complete recall, production scale, real-time operation, or customer savings. PDF reconciliation, scheduling, deployment, and CRM delivery remain outside WS-001.

Reproduce the proof

```
python -m unittest discover -s projects/WS-001-qualified-tender-change-intelligence/tests -v
```